

Temporal-Support Identifiability in Event-Detection Evaluation Pipelines

Anonymous CVPR submission

Paper ID *****

Abstract

001 *An event-detection benchmark identifies each scored sam-*
002 *ple by one nominal timestamp, but a recurrent detector*
003 *forms its prediction from a longer history, and the bench-*
004 *mark does not record how long. We state the missing re-*
005 *quirement as a property of the protocol, temporal-support*
006 *identifiability: the benchmark variables must uniquely de-*
007 *termine the temporal evidence supporting a reported pre-*
008 *dition. This makes effective temporal support part of the*
009 *benchmark specification rather than a hidden property of*
010 *one evaluation implementation. A released SSM-ViT eval-*
011 *uation pipeline provides a direct counterexample. At the*
012 *same evaluated timestamp, with frames, weights, labels and*
013 *evaluator fixed, moving only the chunk boundary changes*
014 *the recurrent history from one to twenty-one windows within*
015 *the pooled score, and changes S5-B by 4.20 mAP over*
016 *all labeled boxes and by 5.23 ± 0.49 mAP on the velocity-*
017 *evaluable subset. Because identifiability is a uniqueness*
018 *property, one controlled pair in which identical recorded*
019 *protocol variables yield different effective support and dif-*
020 *ferent scores is sufficient to show that the protocol under-*
021 *specifies the score. Two supporting diagnostics bound the*
022 *interpretation without being used as recurrent-support es-*
023 *timators: on the predictor side, the newest RVT input win-*
024 *dow has an ablation-sensitivity centroid 23.81 ms before the*
025 *label instant and a nominal-time-matched regression coeffi-*
026 *cient of -2.40 ± 7.18 ms; on the label side, within-exposure*
027 *event timing contains an illumination-locked component.*
028 *Timing comparisons between recurrent detectors therefore*
029 *require the support that produced each score to be stated*
030 *and equalized, a condition an event-detection protocol can*
031 *record with its labels.*

032 1. Introduction

033 A timestamp says when a prediction is scored, not what
034 temporal evidence produced it. Two recurrent detectors
035 evaluated at the same nominal timestamp are therefore not

comparable in time unless the history each one actually used
is also identified.

An event-detection benchmark defines a sample by a
nominal timestamp. A prediction $\hat{s}$ is scored against the
ground truth s^* carried at query time t through a function
 $\rho(\hat{s}, s^*)$ of position, shape and class. Beyond selecting the
label at t , that function contains no variable for the temporal
evidence behind either argument, so any temporal discrep-
ancy enters the score disguised as displacement.

A recurrent detector does not read only the instant the
benchmark names. It computes $\hat{s}$ from an input window and
from state carried across earlier windows, and the label it is
compared against comes from an exposure of finite width,
an annotator or tracker working on the exposed image, and
a geometric transfer into another sensor's frame. The evi-
dence on both sides is an interval, while the protocol records
an instant.

Three terms carry this throughout. The *nominal times-*
stamp is the instant at which the protocol pairs a prediction
with a label. The *effective temporal support* is the past that
can influence the prediction, including state carried by re-
currence. The *label-acquisition support* is the interval over
which the label's evidence was acquired. A nominal times-
stamp determines neither of the other two uniquely, and a
temporal descriptor may be reconstructable from released
artifacts yet remain unreported or heterogeneous within one
pooled score. We call this protocol property temporal-
support identifiability: whether the released protocol and
artifacts uniquely determine the temporal quantities a re-
ported score represents.

We test it with a paired intervention on the released
SSM-ViT evaluation pipeline. In that implementation the
nominally carried state is inert across evaluation chunks,
so the recurrent history available to a prediction is a de-
terministic function of chunk position and ranges from one
to twenty-one windows inside one pooled score. Moving
every chunk boundary 5 windows later re-scores the same
frames under a longer history while frames, weights, ground
truth, evaluator and the evaluated timestamp are held fixed
(3 574 paired boxes, cluster bootstrap over 406 sequences).

Effective temporal support alone changes S5-B by 4.20 mAP over all labeled boxes and 5.23 ± 0.49 mAP on the velocity-evaluable subset (Sec. 3); the second released checkpoint moves by 4.10 and 4.81 points. Every quantity the protocol records stayed fixed, and the score moved. Since temporal-support identifiability is a uniqueness property, this controlled violation is sufficient to show that the nominal timestamp does not uniquely specify the temporal evidence represented by that score.

Contributions.

- *Temporal-support identifiability*: an evaluation property separating the nominal timestamp a protocol records from the effective temporal support and label-acquisition support a reported score represents.
- A paired intervention that isolates recurrent support: the frames, weights, labels, evaluator, and evaluated timestamp are fixed, while only the recurrent history changes. The two audits below support the interpretation and are not second primary results.
- Its measured consequence: support the protocol does not record moves a pooled leaderboard-style score by 4.20 mAP overall and 5.23 mAP on the velocity-evaluable subset, with predictor-side and label-side audits delimiting the result.

Work on event-detection timing takes that index as given and improves what is measured against it: sensor latency is calibrated away [25, 29], compute latency is folded into average precision [9, 30, 35], and the temporal window is redesigned [24, 26]. Each changes what the predictor does at a fixed benchmark index, whereas the object measured here is that index itself and the support a score reported under it pools. The two datasets below instantiate this protocol condition; how often it holds elsewhere is not measured.

Supporting audit: predictor-side support. The event branch examined here [24] consumes a window of 50 ms in ten bins ending on the label instant (Sec. 4.2), so under uniform weighting its evidence is centered 25 ms earlier; ablating each bin on real validation samples puts the ablation-sensitivity centroid at -23.81 ms (Sec. 4.3), the newest window carrying 22.7% of the summed window-ablation sensitivity over a 1000 ms horizon and preceding windows 77.3%. The regression-based temporal coefficient, estimated on nominal-time-matched detections under Eq. 2 with a placebo term and a sign-flip control, is -2.40 ± 7.18 ms, sits 26.2 ms from that centroid and is compatible with zero: it is an association between nominal-time-matched detections, not an estimate of recurrent support, and the nominal timestamp encodes neither (Secs. 4.2–4.5).

Supporting audit: label-acquisition support. Across the eighteen DSEC training sequences whose exposure files resolve, published exposure width varies by a factor of

127 against a frame period fixed at 50 ms, with 32.5 % of frames pinned at an auto-exposure ceiling in six sequences (Fig. 3). Those sequences carry the harmonic series of a 100 Hz mains flicker, and between 95.4 % and 100.0 % of labeled boxes exceed the ninety-ninth percentile of the off-frequency local excess measured on the same events, so within-exposure event timing is not attributable to object timing without separating that component (Secs. 5–5.4).

Identifiability under the metric is then read in the metric’s own units. On the velocity-evaluable subset, displacing ground-truth centers across the measured label-to-centroid interval changes mAP by 0.04 points, and at that displacement the five released checkpoints are not reordered. Separately, over all labeled boxes the nominal validation-split ordering is reproduced in only 42.5 % of sequence-bootstrap replicates. These results place the observed displacement effect at the benchmark’s sampling-resolution floor. Table 2 relates this resolution to the benchmark’s speed composition (Supplement Secs. 11–12).

2. Related work

Sensor and runtime latency. Prior work calibrates sensor latency [25, 29], incorporates runtime latency into average precision [9, 30, 35], or proposes temporal tolerance [36]. These methods address the time from an event to an output and take the benchmark’s temporal index as given. We instead measure that index, testing whether a released benchmark specification uniquely determines the temporal quantities behind its score.

Temporal windows and benchmark identifiability. RVT and SSM-ViT redesign or parameterize temporal windows [24, 26], while adaptive-window methods compare alternative partitions [28]. Event–RGB frequency mismatch and exposure duration have also been noted [22, 23, 37]. The closest accepted work above calibrates sensor or runtime latency or redesigns temporal windows, each holding the benchmark’s temporal index fixed. We make that index the measured quantity, reading the released per-frame windows and exposures and separating predictor support, output-time interpretation, and label-acquisition support.

Benchmark reliability and evaluation analysis. Prior work studies detector calibration, stratified model evaluation, label-induced benchmark instability, and event-vision evaluation protocols [31–34]. These analyses examine uncertainty or error in reported performance, whereas our focus is the temporal support and acquisition semantics that a nominal detection timestamp leaves implicit.

Label time and attribution. Waymo and annotation studies estimate per-object capture or annotation offsets [5–

7]. Event datasets expose per-box time fields, but their degrees of freedom are not necessarily specified. LET-3D-AP and TIDE/DETAD/HOTA decompose or bucket spatial error while taking the label clock as given [1, 11–14].

Event benchmarks and exposure. Gen1, 1 Mpx, DSEC-Det and Ev-3DOD use asynchronous measurements, rectified frames, or interpolated labels [16, 19–21]. LEOD establishes causal input windows [27]. The present work reads the released window and exposure artifacts as protocol evidence and tests whether the resulting temporal quantities are uniquely specified.

3. Primary evidence: the same-frame chunk-boundary intervention

The state the SSM-ViT release carries between evaluation chunks does not affect any of its outputs: it enters as a multiplicative factor on the first coefficient of an associative scan on which the outputs do not depend. Under the released streaming evaluation this bounds the recurrent history available to a detection to the current 21-window chunk, so that history ranges from one window to 21 with the detection’s position in the chunk, and the reported number averages over all 21 positions. Chunk position is a deterministic function of the release’s own index files and is recorded nowhere in the evaluation protocol.

The same-frame intervention. The intervention moves every chunk boundary 5 windows later and changes nothing else: the same frames, the same weights, the same labels, the same evaluator, and the evaluated timestamp of each frame unchanged. 3 574 labeled boxes are scored twice under this shift, once with one to four windows of recurrent history and once with seventeen to twenty, and are compared paired within frame with a cluster bootstrap over the 406 sequences. The long-history arm gains 5.23 ± 0.49 mAP points for S5-B and 4.81 ± 0.52 for S5-S on the velocity-evaluable subset, and over every labeled box the S5-B gain is 4.20 ± 0.36 points. Mean detection confidence, which uses no labels, moves with it by 0.0579 ± 0.0048 and 0.0509 ± 0.0035 . Fig. 1 (bottom) draws one such pair. No quantity the protocol records distinguishes the two arms.

The observational contrast across chunk positions. Reconstructing chunk position over 20 296 labeled frames of the released evaluation and scoring every labeled box, S5-B reaches 34.445 mAP on the frames given one to four windows and 40.226 on those given seventeen to 21 (Fig. 1, top). This contrast compares different frames, so it is differenced against the three RVT checkpoints, evaluated on the same frames with a state that crosses chunk boundaries; that controls for frame-position effects shared with




Figure 1. **Top:** benchmark score against the recurrent history the released streaming evaluation made available, at each of the 21 chunk positions, over every labeled Gen1 box with no displacement applied. Shaded: the one-to-four-window block the release scores together with the rest. Between that block and frames given seventeen or more windows, the two SSM checkpoints rise by 5.78 and 4.83 points; the three RVT checkpoints, scored on the same frames with state that crosses chunk boundaries, rise by at most 0.80. **Bottom:** one frame of the same-frame intervention on S5-B. Changed: the chunk boundary, and with it the recurrent history behind the detection. Fixed: frame, weights, labels (dashed), evaluator and evaluated timestamp. Solid: a detection matched to a label of its class at $\text{IoU} \geq 0.5$; dotted: unmatched; drawn above confidence 0.3 and scored at every confidence. The example is illustrative; over all 3 574 frames the same intervention changes S5-B by 4.20 mAP over all labeled boxes and 5.23 ± 0.49 mAP on the velocity-evaluable subset (Supplement Sec. 13).

them and leaves 5.32 ± 0.53 points for S5-B ($z = 10.0$) and 4.37 ± 0.61 for S5-S ($z = 7.2$), against placebo differences among the RVT checkpoints of at most 0.67 points ($z = 1.56$). Permuting chunk position within each sequence puts both SSM values at least 7.6 standard deviations outside their own null and every placebo inside it. The two readings agree in sign and magnitude, and the paired intervention carries the claim because it holds the frames fixed (Supplement Sec. 13).

What the pooled score reports. One released number therefore pools detections whose recurrent history differs by a factor of 21. Restricting evaluation to the frames receiving 17–21 windows yields an S5-B score 1.75 points above the pooled released-evaluation score. Two systems compared under this protocol at the same nominal timestamps need not have been given the same temporal evidence, and the protocol records no variable by which a reader could tell.

4. Supporting audit: predictor-side temporal descriptors

4.1. Identifiability and the timing estimands

A benchmark releases one tuple per evaluated instant: a nominal timestamp, the labels asserted at it, the predictions scored against them, and the score. Three temporal quantities are latent in that tuple and are not recorded in it: the

temporal support of the predictor’s input and carried state, the effective state time its output refers to, and the acquisition support of the labels. Call such a quantity *protocol-identifiable* when the released protocol and artifacts determine it uniquely. Under this definition, input and exposure intervals are protocol-identifiable, whereas effective predictor timing, recurrent context represented by a pooled score, and within-exposure attribution are not.

Let a predictor P emit a state at query time t . A signed effective-weight model uses bin centers c_k and weights a_k with $\sum_k a_k = 1$, giving the effective timestamp offset $\tau_P = \sum_k a_k c_k$ under constant velocity. The empirical ablation profile instead assigns non-negative sensitivities $s_k \geq 0$ to bins and has centroid $c_s = \sum_k s_k c_k / \sum_k s_k$. These are distinct quantities: zero-ablation sensitivity depends on both the predictor and the input, and is not in general a decomposition of temporal weights. Sections 4.2–4.3 measure this operational input-sensitivity profile over the newest window. Section 4.4 estimates the regression coefficient τ under Eq. 2. Interpreting $-\tau$ as an effective state-time offset requires the absence of residual velocity-proportional spatial bias. Section 4.5 compares $-c_s$ with the regression coefficient τ , equivalently c_s with the regression-implied timestamp $\tau_P = -\tau$.

4.2. The input window and its indexing instant

Two released files fix the input side. The dataset configuration `config/dataset/gen1.yaml` of the detector’s own repository [24] names the representation `stacked_histogram_dt=50_nbins=10`, a window of 50 ms divided into ten bins of 5 ms. Line 405 of `scripts/genx/preprocess_dataset.py` constructs the array `ev_repr_timestamps_us_end` by counting backwards from a label timestamp in steps of the window duration, and extends it between consecutive label timestamps. Each representation is therefore indexed by the end of its window, and those ends are placed on label times. A window attached to a label at time t holds the events of $[t - 50 \text{ ms}, t]$, whose uniform weighting centroid falls 25 ms before the instant t at which the label asserts a state. We read both statements from the released source, and the 25 ms figure is a reference point instead of a property of the trained network.

4.3. The measured influence profile

The newest-window ablation-sensitivity centroid is computed from the measured bin-ablation sensitivities of a released checkpoint. We take 480 samples of the released Gen1 validation split with the released `rvt-t` checkpoint, warm the recurrent state for eight steps, and ablate each bin in turn. For sample i , let Y_i denote the emitted detection tensor from the unablated input and $Y_i^{(-k)}$ the tensor obtained after zeroing temporal bin k while keeping the incoming re-




Figure 2. Input sensitivity and the regression-based temporal coefficient summarize the same nominal timestamp without being equivalent, and are not competing estimates of a common latent timestamp; their point estimates are 26.2 ms apart. Panel (b) shows that average precision on the velocity-evaluable subset is nearly invariant over the measured 23.81 ms label-to-centroid displacement. The 26.2 ms separation is not an estimate of architectural latency. (a) Window-ablation influence of each of the ten bins of the released `rvt-t` input window, over 480 Gen1 validation samples with the recurrent state warmed eight steps, with sample-level standard errors. Error bars are descriptive standard errors across the sampled instances and are not used for inference. Uncertainty of the ablation-sensitivity centroid is evaluated by sequence bootstrap in Supplement Sec. 2. Dashed: the mean over bins. Dotted: the label instant. Solid: the ablation-sensitivity centroid. Dash-dot, with its standard error shaded: the regression-implied timestamp under Eq. 2, $\tau_P = -\tau$, of Sec. 4.4. Gray: the four other checkpoints, each rescaled to its own mean. (b) Average precision against the temporal center displacement δ applied to the ground-truth boxes, by object speed in px/s, each curve relative to its own value at the label instant, with its standard error shaded. The shaded band runs from the centroid to the label instant, spanning 0.61 px at the ninetyeth percentile of label speed against an annotation dispersion of 0.6353 px (Supplement Sec. 3). The mAP span is 0.37 points over all moving boxes and 0.78 points in the 25–50 px/s stratum, where the centroid displacement costs 0.08 points.

current state identical. We define the ablation sensitivity, or influence, of bin k as

$$s_k = \frac{1}{N} \sum_{i=1}^N \frac{\|Y_i^{(-k)} - Y_i\|_2}{\|Y_i\|_2}, \quad N = 480, \quad (1)$$

$$c_s = \frac{\sum_k s_k c_k}{\sum_k s_k},$$

where c_k is the center time of bin k relative to the label instant. The ablated forward pass receives the same incoming recurrent state as the corresponding reference pass, and its returned state is discarded. The profile is a property of the current window and not of a perturbed history.

The profile is balanced between the two halves of the window (Fig. 2a) and yields an ablation-sensitivity centroid of -23.81 ms, compared with -25.00 ms for uni-

form weighting. The released configuration therefore places the measured centroid near the window reference without implying uniform network sensitivity. Across zero, mean, cross-sample replacement, and gradient instruments, the centroid spans -26.35 to -22.35 ms (Supplement Sec. 2). The reported centroid applies to the deterministic 12-sequence subset described in the supplement. This is the current window’s window-ablation sensitivity profile with the recurrent state held fixed, not a decomposition of the full temporal support.

Recurrent support beyond the newest window. Over a 1000 ms horizon, the newest window accounts for 22.7% of the summed window-ablation sensitivity, while the remaining 77.3% lies in preceding windows. The horizon centroid shifts from -161 ms at 500 ms to -299 ms at 1000 ms, so the measured range does not determine a converged full-history centroid. Supplement Sec. 14 characterizes the tail and its fill controls.

Across the released checkpoints and a second architecture. RVT releases three Gen1 checkpoints differing only in embedding width, attention head dimension and FPN depth. SSM-ViT [26] releases two more in which each stage’s ConvLSTM is replaced by an S5 state space layer, at 9.68 and 18.19 M parameters against RVT’s 9.87 and 18.54. It distributes RVT’s own preprocessed Gen1 and inherits its window construction. On the same sequences, and with each driven as its own release specifies, all 5 checkpoints put the newest window’s ablation-sensitivity centroid between -23.76 and -24.72 ms, none further than 1.24 ms, or 25 % of one bin, from the uniform-weight value.

The support within the newest window is not fixed across capacities: over 576 samples of the three RVT capacities, the newest window’s share of a 1000 ms horizon falls from 22.7 % to 19.0 % and the sensitivity centroid moves from -299 to -339 ms.

4.4. Regression-based temporal coefficient

The influence profile is a statement about the network’s input. Even a predictor with uniform local sensitivity can emit a state describing a later instant, and that quantity is measured from the same checkpoint’s own detections.

A detector reporting the state an object occupied τ earlier incurs the physical displacement $-\tau v$, where v is the latent true velocity. We observe only the finite-differenced label velocity $\tilde{v} = v + \eta$ and therefore fit Eq. 2 using $\tilde{v}$. Thus $\tau = -\tau_P$ on the axis of Sec. 4. With $\hat{c}$ the predicted center, c^* the label center, A the area of the ground-truth box, $\hat{u} = \tilde{v}/\|\tilde{v}\|$ and R a quarter turn, b is a two-dimensional intercept, fitted globally or as per-sequence intercepts in the reported specification,

Specification	τ (ms)	placebo ρ (ms)
<i>velocity estimator, no geometric controls</i>		
forward difference	-14.68 ± 4.99	-6.53 ± 1.68
backward difference	$+45.65 \pm 5.59$	-6.15 ± 1.71
centered difference	$+17.51 \pm 5.06$	-7.76 ± 1.92
<i>controls added, centered difference throughout</i>		
+ box side along travel	$+14.62 \pm 5.47$	-7.75 ± 1.91
+ four geometric terms	-7.62 ± 7.01	-4.12 ± 1.90
+ per-sequence bias	-2.40 ± 7.18	-2.66 ± 1.71

Table 1. The regression-based temporal coefficient τ of the released `rvt-t` checkpoint over 22 534 matched boxes of 376 Gen1 validation sequences, and the placebo ρ beside it, the coefficient on the rotated label velocity, which no temporal offset can reach. Standard errors are clustered by sequence. The forward and backward fits above are calibrated in Supplement Sec. 4.

$$\hat{c} - c^* = b - \tau \tilde{v} + \rho R\tilde{v} + \gamma\sqrt{A}\hat{u} + \varepsilon \quad (2)$$

is fitted by stacking both components of every match into one least squares problem, with standard errors clustered by sequence. It is the temporal calibration of [8] carried to a box center, taken as a slope, not as the per-object ratio $-\langle \hat{c} - c^*, \hat{u} \rangle / \|\tilde{v}\|$, whose variance diverges as its denominator approaches zero. For comparison, the per-object ratio evaluated on the same rows has a mean of $+24.22 \pm 27.73$ ms, on the hypothesis evaluated in Sec. 4.5, with an interquartile range of 430 ms. Written this way, Eq. 2 fits the two-dimensional residual directly with $\tilde{v}$ as a regressor. Unlike the per-object ratio, it does not first project the residual onto $\hat{u}$ and divide by $\|\tilde{v}\|$. The reported full specification adds four coordinate/size regressors and per-sequence intercepts to the terms displayed in Eq. 2. The coefficient is a temporal offset only if no velocity-proportional spatial bias remains after these controls. Such a bias is indistinguishable from τ in Eq. 2.

$R\tilde{v}$ is the label velocity rotated a quarter turn. Because a pure temporal offset acts along $\tilde{v}$, not $R\tilde{v}$, ρ serves as a negative-control coefficient. A nonzero value implies residual spatial misspecification. $\sqrt{A}\hat{u}$ absorbs a displacement along travel proportional to box side, which is needed because that residual grows with box size and box size correlates with label speed at $+0.139$ on this split.

The finite difference used for the velocity. $\tilde{v}$ is finite differenced from labels. Label noise enters forward and backward velocity differences with opposite shared-noise signs. The centered difference excludes the current label center and removes this direct same-frame term. Supplement Secs. 1 and 4 derive and calibrate this effect, so the reported regression uses the centered estimator.

Regression dataset and matching. Rows are detections matched to nominal-time ground truth at $\text{IoU} \geq 0.3$ with

centered label velocities available, leaving 22 534 matches from 376 sequences. The estimate is conditional on this matched subset. Supplement Sec. 5 reports association-rule sweeps, including overlap-free matching.

Table 1 is a specification ladder, and the specification was chosen by the placebo, not by τ . Under Eq. 2 with only the box-side term, $\tau = +14.62 \pm 5.47$ ms, but ρ is -7.75 ms at 4.0 standard errors, which indicates a term omitted from Eq. 2 that $R\tilde{v}$ is a proxy for. A split by the sign of the horizontal velocity locates it: ρ changes sign across that split, at -27.99 ms and $+21.84$ ms, which a rotational term would not do. Most label motion here is horizontal, so $R\tilde{v}$ maps a horizontal speed onto a vertical residual. The rotated regressor represents a bias that differs between oncoming and same-direction traffic. Entering box height and image row along $\hat{y}$ and box width and image column along $\hat{x}$, and then a constant vector per sequence, returns the placebo to -2.66 ms at 1.6 standard errors. That is the specification reported, and the criterion that selected it is a coefficient on a regressor orthogonal to $\tilde{v}$.

Under the reported specification, $\tau = -2.40 \pm 7.18$ ms, with a sequence-bootstrap 95% interval of $[-17.0, +11.5]$ ms. Across geometric-control specifications, τ spans -20.2 to $+11.1$ ms. Because the reported specification was selected using the orthogonal placebo coefficient, its nominal standard error does not include specification-selection uncertainty. We treat the point estimate descriptively and use the full specification sweep to assess robustness.

4.5. Relating input-sensitivity timing to the regression-based temporal coefficient

The comparison combines measurements under different recurrent-state conditions. The centroid uses an eight-step warm-up with the incoming state fixed during ablation. The regression uses detections from the validation stream with the released state carried across frames. The 26.2 ms numerical separation is descriptive only: the two quantities are distinct operational summaries and are not competing estimates of a common latent timestamp. Under the reported specification and its annotation-error assumptions, the point-estimate difference is 26.2 ms. The regression coefficient has a standard error of 7.2 ms. This is a regression-coefficient result instead of a latency estimate of this architecture. Expressed on Eq. 2's sign convention, the newest-window centroid corresponds numerically to $\tau = +23.81$ ms. This comparison relates two operational estimands and does not assume that Eq. 2 must recover the ablation centroid. Thus, 23.81 ms is the measured interval between the label instant and the newest-window ablation-sensitivity centroid, whereas 26.2 ms is the point-estimate difference between that centroid and the reported regression coefficient. The measured horizons do not determine a

converged full-history centroid. The variation across specifications is substantial relative to the reported standard error and is not captured by that standard error.

The detector is trained against the ground truth at the label time on a window ending there, so a measured offset near zero is consistent with the training objective indexed at the label time. The objective does not enforce the interval, and the instant the output describes and the centroid of the evidence used to produce it are different quantities that one benchmark timestamp stands for. At the ninetieth percentile of label speed on this split, 25.7 px/s, the 23.81 ms is 0.61 px of displacement, comparable to the annotation dispersion of the labels themselves (0.6353 px), and at the ninety-ninth it is 1.29 px. That ablation-sensitivity centroid and the regression-based temporal coefficient are measured from released artifacts alone.

Effect on benchmark mAP. This experiment is not a test of whether temporal support matters in general. It asks whether Gen1's particular speed composition makes the measured RVT interval visible to its current mAP. We translate each ground-truth box center by $\delta\tilde{v}$, keeping its size and class fixed, a first-order surrogate for temporal displacement under local constant velocity. The released predictions remain unchanged, since a detection has no velocity of its own. Over 20 296 frames, 68.7% of the 40 698 ground-truth boxes carry a centered velocity. Ground-truth boxes without a centered velocity are marked ignored during evaluation, and detections matched to ignored boxes are excluded from TP/FP accounting. The resulting mAP against δ is smooth with its maximum at -10 ms (Fig. 2b). Along the axis of Eq. 2 that is $\tau = +10$ ms, and mAP varies by at most 0.04 points over the interval. This descriptive curve does not by itself identify a preferred temporal value.

The full -60 to $+60$ ms sweep moves mAP by 0.9 points. On the -50 to $+30$ ms grid summarized in Table 2, the all-moving span is 0.37 points. Displacing the ground truth to the window's centroid costs 0.0004 absolute mAP (0.04 percentage points, 0.12% relative to 0.3346). Average precision on the velocity-evaluable subset is nearly invariant over the measured label-to-centroid displacement, and at the common -23.81 ms displacement the ordering of all 5 released checkpoints is unchanged on the validation split. Scoring each checkpoint at its own preferred displacement likewise leaves the validation-split ordering unchanged in the reported strata. Both statements are bounded by the stability of the ordering under sequence resampling at this scale. The smallest timing difference that would reorder an adjacent pair is 43.7 ms over all moving boxes, $1.83\times$ the measured interval, and 11.2 ms in the 25–50 px/s stratum. Separately, on every labeled box instead of the velocity-evaluable subset the displacement acts on, the benchmark's own cluster resampling holds the validation-split ordering

speed (px/s)	share (%)	med. speed	centroid (px)	Δ mAP	span
all moving	100.0	4.0	0.10	0.04	0.37
< 10	72.1	2.1	0.05	0.02	0.11
10–25	19.2	15.0	0.36	0.01	0.33
25–50	7.5	32.3	0.77	0.08	0.78
≥ 50	1.2	57.4	1.37	0.00	0.37

Table 2. The sweep of Fig. 2b by object speed over the 27 943 ground-truth boxes with centered velocity. Other boxes are ignored. Δ mAP is the cost of shifting the ground truth to the newest-window ablation-sensitivity centroid at -23.81 ms. Span is the mAP range over the -50 to $+30$ ms sweep. Both values are reported in points.

in only 42.5 % of replicates, its tightest adjacent pair separated by 0.050 points against a paired standard error of 0.53. A common displacement sweep reveals one localized inversion between `rvt-b` and `S5-S` in the 25–50 px/s stratum. Table 2 locates the flatness in that composition: 72.1 % of the moving boxes travel below 10 px/s, where the displacement is 0.05 px, and the sweep deepens to 0.78 points between 25 and 50 px/s, where it is 0.77 px, past the labels’ own annotation dispersion of 0.6353 px (Supplement Sec. 3). The near invariance therefore characterizes Gen1’s metric resolution, not the absence of a temporal discrepancy.

5. Supporting audit: label-acquisition timing within one exposure

5.1. Exposure and labels

The predictor-side analysis concerns the temporal support of released detector inputs and outputs. DSEC provides a separate exposure-side analysis: published exposure intervals bound the interval within which event timing is observed, while within-exposure event timing contains illumination-locked structure and therefore cannot be uniquely attributed to object timing without separating that structure. DSEC [15] publishes per-frame exposure metadata, and the files resolve for eighteen train sequences. The width varies by a factor of 127 against a frame period fixed at 50 ms, and within one recording by up to a factor of 12.5. The auto-exposure ceiling of $14\,996\,\mu\text{s}$ holds 32.5 % of frames, in six of the eighteen sequences, and the distribution is not continuous over that range (Fig. 3). The gap separates the two exposure regimes, and ceiling exposure forms a distinct stratum.

The released DSEC-Det labels comprise 390 118 boxes on 6693 tracks across 60 sequences. Every label sits on the frame clock, and the median spacing between consecutive label times is $50\,001\,\mu\text{s}$. Over 361 628 triples of consecutive labels on a track, 7.1 % of middle box centers fall within a hundredth of a pixel of their neighbors’ linear interpolant. This rules out exact local linear interpolation as a general description of the released trajectories without identifying

the underlying annotation process. DSEC-Det ships more than one annotation version [18], the supplement names the one measured here. Median box-center image speed is 40 px/s, with 360 px/s at the ninety-ninth percentile.

5.2. Mains-frequency signature and the per-pixel phase test

Swiss mains is 50 Hz, and rectified lighting can produce the tested 100 Hz modulation. The pre-specified 100-Hz line and its harmonics appear in the event-rate spectrum of every ceiling sequence, while the daytime control lacks the same pattern (Fig. 4). Per-pixel phase tests and the pre-specified frequency control establish the same illumination signature.

Per pixel, the Rayleigh statistic $Z = R^2/N$ is $\text{Exp}(1)$ under independent uniform phases. The off-frequency and daytime controls empirically calibrate this nominal null, while the ceiling sequence has median $Z = 2.935$ at 100 Hz and 24.07 % of live pixels reach the nominal $p < 10^{-3}$ threshold under the independent-uniform-phase reference. The illumination attribution therefore relies primarily on the frequency-matched and daytime controls, not on the nominal Rayleigh reference alone.

5.3. All six ceiling sequences, per labeled box

The same statistic taken per labeled box, over the events inside that box during that frame’s own exposure window, extends to every ceiling sequence. Across the six sequences and 16 516 boxes, 95.4–100.0 % satisfy the pre-specified criterion in Supplement Table 2.

5.4. Within-box temporal dispersion under that signature

That signature limits attribution of within-exposure event timing to object timing. The per-box 137-Hz arm is used as a same-event descriptive control, not as an independent spectral-resolution test. The illumination attribution is supported primarily by the long-slice 100-Hz harmonic structure and the daytime control. Within-box dispersion is $206.6\,\mu\text{s}$ against a null of $93.0\,\mu\text{s}$, with an excess of $183.6\,\mu\text{s}$. A frame-level permutation preserving pooled temporal distributions and box counts gives a null of $78.6\,\mu\text{s}$, and the excess remains after the clean-pixel control. Random-position boxes show that the dispersion is not specific to annotated object regions (Supplement Sec. 9), which supports reading the signature as an attribution limit and not as object motion.

6. Scope and limitations

Each supporting audit is conditional on its stated instrument: local window-ablation sensitivity, nominal-time association under the reported regression specification, and a local constant-velocity displacement surrogate. The primary SSM chunk-boundary intervention is established on




Figure 3. Published per-frame exposure width across the eighteen DSEC training sequences whose exposure files resolve: per sequence (left, log axis, median, interquartile range and 5–95 % band, against the fixed 50 ms frame period) and pooled over 21 142 frames (right). 32.5 % of frames sit at the auto-exposure ceiling of 14 996 μ s and zero fall between the largest non-ceiling width of 4207 μ s and that ceiling, separating the two exposure regimes. The six ceiling sequences are those measured in Sec. 5.




Figure 4. Top: 350 μ s slices of the ceiling-exposure sequence `zurich_city_09_a` and of the daytime control, at the peak and trough of the ceiling sequence's own event rate inside its densest published exposure window. Bottom: that rate in 350 μ s bins, each sequence normalized to its own mean. The ceiling rate moves by a factor of 3.0 between its extremes, the daytime rate by 1.1.

unreported degree of freedom of the protocol: it changes the recurrent support behind the reported score while every recorded quantity stays fixed, which is exactly a failure of temporal-support identifiability. Effective temporal support is therefore part of what a recurrent benchmark score means, not an implementation detail.

The result does not single out a canonical temporal-support statistic, nor attribute the shift to one stage of the pipeline. RVT's input sensitivity and its regression-based temporal coefficient remain descriptors rather than estimates of recurrent support, and DSEC's within-exposure timing carries illumination-locked structure that must be separated before it is attributed to object motion.

The actionable rule is short. Two scores at the same nominal timestamp are not performance at the same effective time when the recurrent support behind them differs, so comparisons across recurrent architectures should report or control predictor initialization, reset policy, chunk boundaries, effective history and label-acquisition intervals alongside the score, which makes these quantities protocol-identifiable.

two released SSM-ViT checkpoints under their released streaming evaluator. Temporal-support identifiability is a uniqueness claim, so one controlled pair with identical recorded protocol variables, different effective support, and different score is already a protocol-level counterexample; the second checkpoint tests that the observed violation is not confined to one released weight set. Gen1 and DSEC are separate audits of predictor- and label-side timing, not premises required for that counterexample. All Gen1 quantitative claims are validation-split claims obtained from ported released checkpoints, and sequence-bootstrap claims use the 406 sequences represented in the stored-detection evaluation.

7. Conclusion

In the released SSM-ViT pipeline one pooled score combines predictions with one to twenty-one windows of recurrent history, and moving only the chunk boundary shifts S5-B by 4.20 mAP over all labeled validation boxes and 5.23 mAP on the velocity-evaluable subset, at fixed frames, weights, labels and evaluator. The chunk boundary is an

References

- [1] W.-C. Hung, V. Casser, H. Kretzschmar, J.-J. Hwang, and D. Anguelov. LET-3D-AP: Longitudinal error tolerant 3D average precision for camera-only 3D detection. *arXiv:2206.07705*, 2022. Waymo Open Dataset 3D camera-only detection challenge metric. 3
- [2] S. Tanaka, K. Minoda, F. Watanabe, and T. Horibe. Rethink 3D object detection from physical world. *arXiv:2507.00190*, 2025.
- [3] A. Kugele, T. Pfeil, M. Pfeiffer, and E. Chicca. How many events make an object? Improving single-frame object detection on the 1 Mpx dataset. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2023, pages 3913–3922.
- [4] Z. Zhou, Z. Wu, R. Bouteau, F. Yang, C. Demonceaux, and D. Gin-hac. RGB-Event fusion for moving object detection in autonomous driving. *arXiv:2209.08323*, 2022.
- [5] Waymo Open Dataset. `label.proto`, `field camera_synced_box` and its documentation comment, read 2026. 2
- [6] P. Sun et al. Scalability in perception for autonomous driving: Waymo Open Dataset. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [7] A. J. Miro, L. af Klinteberg, B. Timus, A. Asefaw, A. Khoche, T. Gustafsson, S. S. Mansouri, and M. Daneshmand. Correcting and quantifying systematic errors in 3D box annotations for autonomous driving. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2026. 3
- [8] T. Qin and S. Shen. Online temporal calibration for monocular visual-inertial systems. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2018. 5
- [9] M. Li, Y.-X. Wang, and D. Ramanan. Towards streaming perception. In *European Conference on Computer Vision (ECCV)*, 2020. 2
- [10] H. Mao, X. Yang, and W. J. Dally. A delay metric for video object detection: What average precision fails to tell. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019, pages 573–582.
- [11] D. Bolya, S. Foley, J. Hays, and J. Hoffman. TIDE: A general toolbox for identifying object detection errors. In *European Conference on Computer Vision (ECCV)*, 2020. 3
- [12] D. Hoiem, Y. Chodpathumwan, and Q. Dai. Diagnosing error in object detectors. In *European Conference on Computer Vision (ECCV)*, 2012.
- [13] H. Alwassel, F. Caba Heilbron, V. Escorcia, and B. Ghanem. Diagnosing error in temporal action detectors. In *European Conference on Computer Vision (ECCV)*, 2018.
- [14] J. Luiten, A. Ošep, P. Dendorfer, P. Torr, A. Geiger, L. Leal-Taixé, and B. Leibe. HOTA: A higher order metric for evaluating multi-object tracking. *International Journal of Computer Vision*, 2021. 3
- [15] M. Gehrig, W. Aarents, D. Gehrig, and D. Scaramuzza. DSEC: A stereo event camera dataset for driving scenarios. *IEEE Robotics and Automation Letters*, 2021. 7
- [16] D. Gehrig and D. Scaramuzza. Low-latency automotive vision with event cameras. *Nature*, 629:1034–1040, 2024. 3
- [17] J. Pang, L. Qiu, X. Li, H. Chen, Q. Li, T. Darrell, and F. Yu. Quasi-dense similarity learning for multiple object tracking. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.
- [18] K. Wang, S. Liu, H. Shi, L. Shi, and H. Chen. Beyond duality: A hybrid framework of leveraging shared and private features for RGB-Event object detection. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. 7
- [19] P. de Tournemire, D. Nitti, E. Perot, D. Migliore, and A. Sironi. A large scale event-based detection dataset for automotive. *arXiv:2001.08499*, 2020. 3
- [20] E. Perot, P. de Tournemire, D. Nitti, J. Masci, and A. Sironi. Learning to detect objects with a 1 megapixel event camera. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [21] H. Cho, J.-Y. Kang, Y. Kim, and K.-J. Yoon. Ev-3DOD: Pushing the temporal boundaries of 3D object detection with event cameras. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. 3
- [22] L. Sun, C. Sakaridis, J. Liang, Q. Jiang, K. Yang, P. Sun, Y. Ye, K. Wang, and L. Van Gool. Event-based fusion for motion deblurring with cross-modal attention. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 2
- [23] T. Kim, J. Lee, L. Wang, and K.-J. Yoon. Event-guided deblurring of unknown exposure time videos. In *European Conference on Computer Vision (ECCV)*, 2022. 2
- [24] M. Gehrig and D. Scaramuzza. Recurrent vision transformers for object detection with event cameras. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 2, 4
- [25] H. Liu, S. Peng, L. Zhu, Y. Chang, H. Zhou, and L. Yan. Seeing Motion at Nighttime with an Event Camera. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 2
- [26] N. Zubić, M. Gehrig, and D. Scaramuzza. State space models for event cameras. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 2, 5
- [27] Z. Wu, M. Gehrig, Q. Lyu, X. Liu, and I. Gilitschenski. LEOD: Label-efficient object detection for event cameras. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16933–16943, 2024. 3
- [28] Z. Sui, H. Hao, W. He, S.-H. Lee, and W. Wang. Adaptive spatial-temporal window: unlocking the potential of event cameras in heterogeneous velocity scenarios. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 946–955, 2026. 2
- [29] Y. Yang, J. Liang, B. Yu, Y. Chen, J. S. Ren, and B. Shi. Latency correction for event-guided deblurring and frame interpolation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 24977–24986, 2024. 2
- [30] J. Chu, R. Zhang, C. Yang, Z. Yu, Z. Bu, H. Liu, F. Röhrbein, A. Knoll, G. Chen, and C. Jiang. Bridging the latency gap with a continuous stream evaluation framework in event-driven perception. *Nature Communications*, 17:2441, 2026. 2
- [31] S. Kuzucu, K. Oksuz, J. Sadeghi, and P. K. Dokania. On calibration of object detectors: Pitfalls, evaluation and baselines. In *European Conference on Computer Vision (ECCV)*, 2024. 2
- [32] R. Fogliato, P. Patil, M. Monfort, and P. Perona. A framework for efficient model evaluation through stratification, sampling, and estimation. In *European Conference on Computer Vision (ECCV)*, 2024. 733
- [33] C. G. Northcutt, A. Athalye, and J. Mueller. Pervasive label errors in test sets destabilize machine learning benchmarks. In *NeurIPS Datasets and Benchmarks*, 2021. 736
- [34] B. Ercan, O. Eker, A. Erdem, and E. Erdem. EVREAL: Towards a comprehensive benchmark and analysis suite for event-based video reconstruction. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2023. 2
- [35] J. Yang, S. Liu, Z. Li, X. Li, and J. Sun. Real-time object detection for streaming perception. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. 2
- [36] R. Salles et al. SoftED: Metrics for soft evaluation of time series event detection. *Computers & Industrial Engineering*, 2024. 2
- [37] Haitian Zhang et al. Frequency-adaptive low-latency object detection using events and frames. *arXiv:2412.04149*, 2024. Unrefereed preprint. 2